

SUNIL BHATT

Data Scientist | AI Engineer

☎ 7895295994 ✉ sbsonu7573@gmail.com 🌐 LinkedIn 🐙 GitHub 📍 Pune, India

PROFESSIONAL SUMMARY

Data Scientist having 4+ years of experience building production ML and LLM systems, specializing in RAG, MLOps, and LLMOps (prompt engineering, evaluation, guardrails, monitoring) to deliver scalable, reliable AI solutions.

WORK EXPERIENCES

Data Scientist | Full-time

Apr 2024 - Present

Cognizant

On-site • Pune, India

- Engineered production-grade RAG and Agentic AI pipelines for 500+ page documents, improving retrieval accuracy by ~35% and reducing hallucinations by ~40% using RAGAS.
- Designed LLM-powered document intelligence and autonomous workflow systems, reducing manual analysis time by 60–70% across enterprise workflows.
- Built multi-agent AI systems with tool-calling, memory, and retrieval orchestration for automated document analysis and decision-support tasks.
- Developed LLM monitoring and observability pipelines to track response quality, latency, hallucinations, and drift for continuous optimization.
- Reduced infrastructure costs by ~25% through optimized batching, caching, vector retrieval optimization, and API utilization strategies.
- Implemented CI/CD pipelines using Docker and GitHub Actions with automated drift monitoring and auto-retraining, enabling zero-downtime deployments and improving model stability by ~20%.

Tech Stack : LLMops, RAGAS, XGBoost, FastAPI, Celery, Redis, Docker, GitHub Actions

Junior Data Scientist | Full-time

Jun 2022 - Mar 2024

Cognizant Technology Solutions

On-site • Pune, India

- Developed customer segmentation model improving classification accuracy by ~18% across 10 risk groups.
- Increased KS score from 28.5 → 30 (+5% lift) improving credit risk separation capability.
- Automated reporting pipelines reducing manual effort by 5–7 hours/week (~30% productivity gain).
- Improved data pipeline efficiency by ~20% through optimized SQL queries and processing logic.

Tech Stack : Python, XGBoost, Pandas, NumPy, Scikit-learn

PROJECTS

IntelliDeal AI — M&A Intelligence Platform

Present

AI Engineer

- Developed a **RAG pipeline** for processing **500+** page financial documents, achieving a **0.89 faithfulness score** with a response time of **<10s**.
- Integrated a **LLMOps workflow** including prompt tuning, evaluation (RAGAS), and response validation to ensure production-grade reliability of LLMs.
- Created a **hybrid recommendation engine** for over **130** companies, enhancing match relevance by approximately **30%**.
- Trained a **deal prediction model** achieving **87% accuracy** with explainability using SHAP.
- Designed an **asynchronous ingestion system** capable of handling **100+** concurrent documents with a zero failure rate.
- Reduced analysts' research time by approximately **65%**.

Tech Stack : LangChain, LangGraph, Pinecone, LLMops (Prompt Engineering, RAGAS)

PROJECTS

MultiDoc AI Chatbot — Multi-Document Intelligence |  Github

Feb 2026 - Mar 2026

Individual

- Developed a conversational AI agent that processes multiple simultaneous document uploads, delivering natural language answers with source citation and page attribution.
- Enabled cross-document retrieval across multiple files simultaneously, improving answer accuracy by ~30%.
- Reduced response time by ~40% using optimized embedding + retrieval pipeline.
- Built conversation memory improving multi-turn query relevance by ~35%.
- Created a document comparison feature that identifies contradictions, similarities, and gaps across uploaded documents on demand.

Tech Stack : LangChain, LangGraph, Pinecone, LLMOps (Prompt Engineering, RAGAS), FAISS

Production-Grade Vehicle Insurance Prediction System

Data Scientist

- **Developed** an end-to-end **ML pipeline**, reducing manual effort by **approximately 70%**.
- **Accelerated** model deployment time from days to minutes, achieving **~90% faster** results.
- **Implemented** drift detection, maintaining model accuracy within **±3%** variance in production.
- **Decreased** retraining costs by **~25%** through automated pipelines.

Tech Stack : XGBoost, Scikit-learn, MLflow, SHAP, FastAPI, Docker, GitHub Actions, CI/CD

SKILLS

Databases :	SQLite, MySQL, PostgreSQL, MongoDB
Frameworks & Libraries :	Scikit-learn, XGBoost, Pandas, NumPy, TensorFlow, PyTorch, FastAPI, Flask
Programming Languages & Ai Llm	Python, SQL, AI/LLM: LangChain, LangGraph, LlamaIndex, CrewAI, Agentic AI, RAG, Vector Databases
Tools Platforms :	MLflow, Docker, CI/CD, AWS EC2, ML Pipelines, Model Deployment (API, Batch), Git, GitLab, GitHub, Kubernetes, Jenkins, GitHub Actions, Linux

EDUCATIONS

BTKIT Dwarahat | Almora

B.Tech in Mechanical Engineering

Aug 2015 - Jun 2019

Percentage : 70

CERTIFICATIONS

Introduction to Artificial Intelligence -WorldQuant University
Data Science & Machine Learning -The Skillians